

NEKETHA SURESH

US Citizen | San Francisco | [LinkedIn](#) | [GitHub](#) | +1 (704) 222-6867 | sureshneketha@gmail.com

SUMMARY

Full stack engineer who has often been the early engineer on small teams, the person who builds the product and the pipeline that ships it. Over three years has taken clinical and B2B tools from rough prototype to production, owning React and FastAPI end to end, standing up the containerized CI/CD that deploys them, and turning real user behavior into the signal that makes the product better. Drawn to the developer experience work that makes a whole team faster: design systems, review tooling, and AI coding assistants. Learns a new stack by shipping on it.

EXPERIENCE

BlueRobins *Founding Engineer (Full Stack, Contract)*

Berkeley, CA

UC Berkeley SkyDeck Pad-13 · Dec 2025 – Apr 2026

- Rebuilt an early stage codebase into a production Next.js and PostgreSQL platform, shipping the first live deployment in 5 weeks and leaving a stable v1, including AI powered chat, in active customer use.
- Designed a multi role PostgreSQL schema with RBAC across 4 user types spanning enrollment, scheduling, and payments.
- Integrated Stripe webhook handling for 7 subscription lifecycle events (failed payments, plan changes, refunds), powering reliable subscription commerce.
- Accelerated prototype to production with GitHub Copilot and Cursor, and onboarded 3 interns through 30+ PR reviews, written code quality standards, and a structured ramp to independent contributors.

Sedna Consulting Group *Artificial Intelligence Intern*

Remote (US)

May 2025 – Aug 2025

- Partnered with internal stakeholders to scope and prototype a cloud based multi agent RAG system for RFP review (Python, FastAPI, LangChain) with retrieval and drafting agents, a human in the loop, and source documents stored in Amazon S3, iterating across 3 review cycles with 6 reviewers and cutting manual review time about 45% on covered workflows.
- Owned the feature end to end, from requirements through post launch tuning, containerizing the service with Docker and shipping it to the team's AWS EKS cluster, and turning a rough concept into a shipped internal product.
- Designed, built and demoed a real time offline LLM application (Ollama, Svelte) for the NACo Conference.

Apollo Trident *Full Stack Software Engineer*

Remote (US)

May 2022 – May 2025

- Embedded as an early engineer on a clinical triage platform used daily by 150+ doctors, nurses, and intake staff across 6 hospital departments, wearing full stack and infrastructure hats and growing from frontend tickets to owning mid sized features and the delivery pipeline.
- Contributed to building full stack web applications across Python, TypeScript, JavaScript, React, HTML, and CSS: intake UI, FastAPI services, a PostgreSQL schema spanning 50+ symptom and vitals fields, SQLite for local and test data, and Amazon S3 for clinical document and image storage.
- Helped build and operate the platform's GitLab CI/CD and containerized deployment, shipping Docker images to Kubernetes on AWS EKS with Helm charts, which cut release friction across the team.
- Supported the build and maintenance of a shared frontend component library and design tokens used across product modules, keeping interfaces consistent and accelerating UI build.
- Built internal dashboards that pulled live data from the platform's APIs into SQL backed views, giving the team visibility into adoption, usage, and release health.
- Built and owned an override with reason flow on the triage recommendation engine that turned live clinician usage into an evaluation signal, lifting reason capture from near zero to 90%+ and recommendation acceptance to about 80%, and cut average intake from about 12 minutes to under 4.
- Worked in an Agile team, integrated GitHub Copilot to speed development, ran user testing and a 3 week pilot per release, and mentored newer engineers.

FEATURED PROJECT

Agentic Ingredient Analyzer (2026): Consumer web app (Next.js, FastAPI, Python) that breaks down product ingredient lists with a multi step LLM agent: tool use to look up ingredients, RAG over a curated reference corpus (Postgres and pgvector), and structured outputs that flag concerns and explain them in plain language, gated by an evals harness that catches quality regressions before release.

SKILLS

Languages & Frontend: Python, TypeScript, JavaScript (ES6+), SQL, HTML5, CSS3, React, Next.js, Tailwind CSS; design systems, CSS tokens, shared component libraries

Backend & APIs: FastAPI, Node.js, REST APIs, webhooks, PostgreSQL, SQLite, Supabase, Redis, Amazon S3

DevOps & Infra: Docker, Kubernetes, AWS EKS, Helm, GitLab CI/CD, GitHub Actions, CI/CD, git workflow automation, data pipelines, Agile

AI / LLM Tooling: Prompt engineering, AI agents and tool use, agent development, AI coding assistants (GitHub Copilot, Cursor), RAG, vector search (pgvector), evals, Langfuse observability, structured outputs; OpenAI, Anthropic (Claude), Gemini; LangChain

EDUCATION

University of North Carolina at Charlotte M.S. Computer Science, GPA 3.77 / 4.00

2023 – 2025

Vellore Institute of Technology B.S. Computer Science, GPA 8.2 / 10

2019 – 2023